

Service Capacity and On-Call Rate Card 2026

Engineering Operations - reviewed quarterly - effective 1 July 2026

Table 1. Funded capacity by service

Service	Funded RPS	p95 target	Scale trigger
checkout-api	1,200	180 ms	sustained 70% for 2 weeks
payment-worker	400	900 ms	queue depth above 5,000
search-api	2,500	250 ms	cluster CPU above 65%
catalogue-api	3,000	120 ms	sustained 75% for 2 weeks
notify-worker	150	2,000 ms	queue depth above 20,000
reporting-batch	n/a	n/a	run time above 90 minutes

This rate card records the capacity limits each Engineering service is funded for in 2026, the thresholds at which it must be scaled, and the on-call cover that applies outside working hours. It is reviewed every quarter by the Engineering Manager together with Finance. Where a service exceeds its funded ceiling for two consecutive weeks, the owning team raises a capacity request rather than scaling silently.

Capacity figures are steady-state ceilings, not peak survivability: a service may briefly exceed its ceiling during an incident or a planned migration. Sustained breach is what the review looks at. The p95 latency target in the table is a service objective, not an alerting threshold; alerting fires at 1.5 times the target for ten consecutive minutes.

On-call cover is organised in two tiers. Tier 1 is the owning team's own rota and answers first for its services. Tier 2 is the platform rota and is paged only when Tier 1 does not acknowledge within the stated window, or when an incident spans more than one service. Paging Tier 2 directly is permitted for SEV1 incidents only.

Cost recovery is internal: each team is charged the monthly rate shown for the tier it subscribes to, and the charge does not change with the number of pages received. A team that cancels Tier 2 cover keeps Tier 1 and loses cross-service escalation.

Table 2. On-call tiers and charges

Tier	Hours	Acknowledge within	Monthly rate (MYR)
Tier 1 - owning team	18:00-08:00 and weekends	15 minutes	4,800
Tier 2 - platform	24/7	30 minutes	9,600
Tier 2 - platform (SEV1)	24/7	10 minutes	included in Tier 2

Escalation to the database on-call engineer is charged to the platform rota, not to the owning team, because the database tier is shared. A team may request a temporary uplift of its funded RPS for a launch; uplifts are approved for at most 30 days and revert automatically.

Reporting-batch has no funded RPS because it is a scheduled job rather than a request service. Its capacity is measured by run time: a run exceeding 90 minutes triggers a review of the query plan before any additional workers are funded.

Services not listed here are unfunded and run on best-effort shared capacity. An unfunded service may not declare a SEV1: it must first be added to this rate card at the quarterly review.